

The mass gap: current position of this programme

This is the synthesis of the notes written since the goal was set to a proof of the Yang–Mills existence and mass-gap conjecture. It states what has been proved here, what has been imported, what has been ruled out, and what remains, with constants explicit and each claim labelled.

The short version: the conjecture decomposes into six named statements, of which two are proved here and one is imported. The clause $m < \infty$ is closed as far as variational methods reach, and reduced to the existence of the theory plus the nontriviality of one flowed correlator. The clause $m > 0$ is untouched, and **seven routes to it have been closed by explicit computation**, in two families: those built from operator norms and spectra, and those built from the ground state of the Hamiltonian. Both families fail for reasons now stated exactly, and the second failure identifies what the Euclidean formulation supplies that the Hamiltonian one does not.

1. The decomposition

The obligations map turns “prove the mass gap” into six statements about the Kogut–Susskind Hamiltonian

$$H = \frac{\hbar c}{a} \left[\frac{g^2}{2} \sum_{\ell} (-\Delta_{\ell}) + \frac{2}{g^2} \sum_p (N - \text{Re tr } U_p) \right], \quad \Delta_{a,L} = \frac{\hbar c}{a} \delta(g; N_s, G).$$

	statement	status
T1	finite lattice: unique physical ground state, $\delta > 0$	proved here
T2	$\delta \geq \gamma(g^2/2)C_2$ for $g \geq g_0$, uniform in volume	proved here from an imported theorem; explicit for the Wilson transfer matrix, $g^2 \geq 176$
T2'	$\liminf_{N_s} \delta > 0$ for every $g > 0$	open; false for $U(1)$, and for $SU(N \geq 5)$ with the Wilson action (numerical evidence)
T3	$\delta_{\infty}(g)/(a\Lambda_{\text{lat}}(g)) \rightarrow m/(\hbar c\Lambda) \in (0, \infty)$	open
T4	continuum infinite-volume theory with the axioms	finite-volume ultraviolet stability known
S	small volume: $\Delta = \delta_1 g^{2/3} \hbar c / L [1 + O(g^{2/3})]$	upper side proved, lower side fixed-lattice only

The conjecture is T3, given T4, with a gap $\delta_{\infty} > 0$ on an interval $(0, g_1)$ along the scaling curve; T2' is the stronger statement that the lattice action has a gap at every coupling, which fails for $SU(N \geq 5)$ with the Wilson action (precision of 2026-09-23; the openings note records four ways in). Taken as a lattice statement, the critical-coupling note shows that T2' is exactly the absence of a zero-temperature bulk phase transition, hence on its second-order branch the uniqueness of the continuum limit.

2. What is proved here

T1 (obligations map §2): compact resolvent, a simple strictly positive ground state, a positive gap inherited by the physical sector, for every a , N_s , g and compact connected G .

T2 in Hamiltonian form (T2 note): the electric term is classical in the Peter–Weyl partition and the magnetic term is a bounded range- $\{0, 1\}^3$ perturbation with $\beta = 48N^2/[g^4(N^2-1)]$, so Yarotsky’s theorem gives $\Delta_{a,L} \geq \gamma(g^2/2)C_2\hbar c/a$ for $g \geq g_0$, uniformly in the lattice size, with exponential clustering.

A Lieb–Robinson bound (note): velocity $v \leq 32eNc/g^2$, uniform in volume, because the unbounded part of H is a sum of commuting single-link Laplacians; hence a strong-coupling correlation length $\xi \leq C'aN/(\gamma C_2 g^4)$.

Upper bounds (note, Polyakov note): the Feynman–Bijl inequality with explicit constants; the strong-coupling gap pinned to order $g^2\hbar c/a$ on both sides; the abelian gap bounded by the plaquette-sine structure factor; the proof that **no operator linear in the electric field with c-number coefficients is gauge invariant for a non-abelian group**; and a necessary condition on the Polyakov-loop susceptibility.

The finiteness clause reduced (moment hierarchy): $m \leq M_{k+1}/M_k$ for every k , with the ratios decreasing, so the f -sum rule is the weakest member and the next one, $m \leq \hbar C_0(0)/\int_0^\infty C_0$, needs only T4 and $C_0 \not\equiv 0$. The free-field value of the weakest member is $4.26\hbar c/\sqrt{8t}$, computed in closed form and confirmed by an independent sum rule (free-field note).

Exact structural identities (note): $\Delta V = 4C_2(N|\mathcal{P}| - V)$, so the magnetic energy shifted by its Haar mean is a Laplacian eigenfunction; $|\nabla V|^2 \leq 16V$ for $SU(2)$; and an exact ground-state sum rule tying $\text{Var}_\Omega(V)$ to kinetic quantities.

Flow facts (conjugation, lattice truncation): the Wilson flow is a diffeomorphism of $G^\mathcal{E}$, unitarily implemented, so conjugation preserves the spectrum; the lattice flow carries no coupling and a doubling is dimensionless flow time $1/2$, so truncating the conjugated Hamiltonian at range Ka costs Ce^{-cK} with pure constants.

3. Seven closed routes

route	why it closes
variational upper bounds, for $m > 0$	a spectral measure with all negative moments finite can have $m = 0$
expansion around the free theory	the bound reaches m only at the confinement scale, where its coefficient is nonperturbative
blocking by projection	the boundary grows like M^2 while the block gap does not; the criterion worsens by $3M^2/8$
conjugation by the flow	the spectrum is invariant, so the gain is the truncation, which is cheap
flow before decimation	sup norms are bijection-invariant and spectra conjugation-invariant
Agmon estimate on the ground state	$\ \nabla V\ _\infty$ is extensive, so it bounds the global excess only
pointwise Gibbs domination	the comparison function bites only above the extensive threshold $V_* \simeq N \mathcal{P} $

The first five exhaust the methods built from operator norms and spectra; the last two exhaust the standard ways of extracting a local statement from the ground-state eigenvalue equation. Each closure is a computation with explicit constants.

The reason the last two fail, stated once. Every inequality derived from $-A'\Delta\Omega + (B'V - e_0)\Omega = 0$ compares the total potential with the total energy, and e_0 is extensive, so it separates configurations only by their global excess. The Euclidean weight $e^{-S_w} = \prod_p e^{-(2/g_E^2)V_p}$ factorizes over plaquettes and separates them locally by construction. That is the structural advantage of the Euclidean formulation, derived here rather than assumed.

4. Where the non-abelian structure has been isolated

Three statements distinguish the groups, and any proof of T2' must use one of them.

1. *The commutator potential* of the zero-momentum sector is identically zero for an abelian group; where present, the transverse zero-point energy confines the valley and produces the gap $\delta_1 \hbar^{4/3} g^{2/3} m^{-2/3}$ (G07, claim C133).
2. *Gauge invariance blocks the photon channel*: no c-number-coefficient operator linear in E commutes with Gauss's law for a non-abelian group, so the excitation through which the abelian gap closes must carry a Wilson line (upper-bound note Proposition 6).
3. *The one-loop valley potential* $U(a) = \frac{2\hbar c}{\pi^2 L} \Phi(La)$ vanishes identically in the abelian theory, so the holonomy is free (valley note §3b).

5. What remains

obligation	kind
a gap on $(0, g_1)$ along the scaling curve	the conjecture's lattice half; T2', closedness of the gapped set, is sufficient
openness of the gapped set	technical: gap stability without frustration-freeness
T3	the conjecture, given that gap
T4	construction
local large-field control	available via the transfer-matrix identification: $ \Omega ^2$ is the Euclidean time-slice marginal
the large-field region itself	closed in measure (cost $\eta^2/(4g^2)$ sharp, entropy $\eta^2/(8\pi^2)$); the Hamiltonian chain needs an operator inequality there, whose proof loses $C_V \partial N \hbar c/a$ on the small-field side, a negative perturbation that binds; the polymer expansion consumes measure statements instead

The pattern across every note is uniform: **statements uniform in the cutoff are renormalization statements, and statements local in space need a local weight.** The strong-coupling theorem is uniform in the volume at fixed cutoff; the small-volume theorem is fixed-lattice; the flowed bounds are finite at fixed lattice and need the flow's renormalization in the continuum; the ground-state estimates are global because the eigenvalue equation is.

6. Assessment

The route that survives is the constructive one: conjugate by the flow, which is exact; truncate, which is cheap, inside the large-field region too once the range grows like η ; decimate, which is the problem; and iterate about twenty-one times for $SU(3)$ from $g_{UV}^2 = 1/2$ ($SU(3)$ constants §3). The identification of Ω^2 with the Euclidean time-slice marginal (transfer note) imports every local Euclidean estimate, and the large-field region is then closed in measure: cost $\eta^2/(4g^2)$, sharp; entropy $\eta^2/(8\pi^2)$; the flow's growth there is the Nielsen–Olesen mode and cannot be improved. What the Hamiltonian chain of lower bounds consumes from that region is an operator inequality, whose proof loses a negative term of order $\hbar c/a$ on the small-field side, of exactly the kind that binds (operator-inequality note). The Euclidean polymer expansion consumes the measure statements directly, which is the reason the constructive programme is Euclidean, now stated as a computation.

The coupling map for $SU(3)$. Three regions, with numbers:

region	statement	source
strong, $g^2 > 444$ (Dobrushin) or $g^2 \geq 1056$ (polymer)	$\Delta_W \geq (\hbar c/a) \log(g^2/432)$, resp. $4 \log(g^2/1056)$, uniform in volume, Wilson transfer matrix	Dobrushin note, polymer note
strong, Kogut–Susskind: explicit threshold withdrawn 2026-10-02 (formerly $g^2 \geq 388$, 79 adjacent)	finite-volume bound $\frac{8}{3}g^2 - 12 P /g^2$ only; the volume-uniform KS gap rests on Yarotsky’s existential threshold	continuous-time note
weak, $1/g^2 \gtrsim 10^2$	small-field expansion applies, crude window $C_1 g \leq \eta \leq C_2$	operator-inequality note §4
intermediate, $10^{-2} \lesssim g^2 \lesssim 4 \times 10^2$	no expansion applies; a finite-volume mixing condition (Dobrushin–Shlosman) would give the gap coupling by coupling	finite-verification note

The strong boundary is explicit and the weak one is crude; the width of the intermediate region in one-loop doublings is set almost entirely by the weak side, about 10^3 with the present window and about 20 if the small-field expansion reaches $g^2 \sim 1/2$ (threshold note §4). The Kogut–Susskind threshold from Yarotsky’s theorem, made explicit, is $g_0^2 \sim 10^{101}$ and plays no role in the map.

The map as a theorem. The conditional theorem states the whole map as two hypotheses, control of the blocking steps from the weak side to the coupling where $\xi \simeq a$ (H1) and certified mixing at that one coupling on one box (H2), with the conclusion $m \geq \hbar c \min(\gamma_*, \gamma')/a_*$, where $a_* \simeq 0.1$ fm is the physical spacing at the verified coupling. Progress is progress on H1 or H2.

The two reasons to stop, researched. A Griffiths-type inequality would transfer control from weaker to stronger coupling only and would replace the strong-side steps, never the weak side (research note); a certified verification at $\beta_W \simeq 6$ is beyond every certified method. On the weak side, one explicit small-field step (part 1, part 1b) has block Poincaré constant $\frac{4}{9}$, fluctuation size $\frac{3}{2}g$ per link, a proved propagator decay rate of 10^{-3} per lattice unit against an order-one truth, and a remainder threshold scaling as κ^4 that lies at $g^2 \sim 10^{-12}$ even for an ideal rate: **twelve orders of magnitude below the physical onset of the running at $g^2 \simeq 1$** , where the strong side’s deficit is a factor 400. Constant-chasing cannot close the weak side; H1 is a methods problem at order-one coupling.

Division of labour. Hamiltonian methods for T1, T2, the small-volume theorem, the upper bounds, the exact identities and the final gap extraction; the Euclidean polymer expansion for the renormalization steps; and, for the intermediate region, the finite-volume mixing conditions, which reduce the gap at a given coupling to a finite verification and turn the problem into the entry of the trajectory of effective interactions into the open set of completely analytical interactions, a sufficient route whose two-number form is the meeting of the reach g_{RG}^2 of the small-field renormalization with the reach g_{DS}^2 of the verification (finite-verification note).

What this programme has added is a map with constants: six named statements, two proved, the finiteness clause reduced to a single correlator, seven routes closed with explicit numbers, the large-field region closed in measure and its obstruction in norm identified, three places where the non-abelian structure is isolated in a usable form, the exact identities and sum rule of Section 2, and explicit thresholds on both sides of the intermediate region. Several of the closures are

corrections of claims made earlier in the same series, and each is recorded with the computation that forced it.

7. Where the map ends

The mass gap for $SU(3)$ in four dimensions is not proved here. Both sides of the confinement scale carry explicit constants, the two standing reasons for stopping have been researched rather than left as reasons, and what remains can be stated in one sentence.

The two reasons, researched. A Griffiths-type correlation inequality would give monotonicity of the string tension and of decay rates in channels with vanishing mean, and it transfers control from weaker to stronger coupling only; on this map it replaces the three strong-side blocking steps and leaves the weak side untouched, so it is a simplification and not an unblocking. Even the vacuum-sector gap's monotonicity would not follow, the truncated plaquette correlator having two terms that both increase. A certified verification of the mixing condition at $\beta_W \simeq 6$ is a supremum over boundary conditions of an integral in thousands of dimensions, while certified polymer enumeration converges only where the expansion does, at β_W of order one; so it cannot be supplied where it is needed, by any method known to this author.

The asymmetry of the two sides. The strong side is proved to $\beta_W = 0.0135$ against a physical crossover at 5.7, a factor of 400. The weak side, computed explicitly for one blocking step, has its threshold controlled by the fourth power of the fluctuation propagator's decay rate, which no source states explicitly and which two explicit arguments put at 10^{-3} per lattice unit against an order-one truth; the threshold lands at $g^2 \sim 10^{-12}$ even with an ideal rate, twelve orders of magnitude below the physical onset of the running at $g^2 \simeq 1$. Sharpening constants within the present method closes neither side, and the weak side by a margin that no sharpening addresses.

What is open. Control of the renormalization steps at couplings where the fluctuation is not small compared with the nonlinearity, with no expansion in any parameter. That is the mass-gap problem for $SU(3)$, stated as precisely as this programme can state it, and it is new mathematics.